

III. Systemcall implementation

Brk() systemcall implementation

Implementing the `brk()` systemcall involves modifying the end of a process's data segment (i.e., the heap). Here's a high-level overview followed by a basic example of how this could be done in a Linux-like kernel environment (for learning or OS development purposes).

Overview of `brk()`

Purpose: Set the end of the data segment (heap) to a specified value.

Signature: `int brk(void* end_data_segment);`

If successful, it returns 0; on error, it returns -1 and sets `errno`.

Kernel-Level Pseudocode for `brk()` Assume:

current is the current process descriptor.

The process has a `mm_struct` that holds memory segment info.

C code

```
SYSCALL_DEFINE1 (brk, unsigned long, end_data_segment)
{
    struct mm_struct *mm = current->mm; unsigned long old_brk, new_brk;

    down_write(&mm->mmap_sem); // lock

    old_brk = mm->brk; new_brk = PAGE_ALIGN(end_data_segment);

    if (new_brk < mm->start_brk || new_brk > mm->end_code) { up_write(&mm
->mmap_sem); // unlock return -ENOMEM;
    }
}
```

```

if(new_brk>old_brk){
    if(expand_heap(mm,old_brk,new_brk)<0){
        up_write(&mm->mmap_sem); return-ENOMEM;
    }
}

```

```

mm->brk=new_brk; up_write(&mm->mmap_sem);

```

```

return0;
}

```

User-SpaceExample(forunderstanding) ccode

```

#include<unistd.h> #include<stdio.h>

```

```

intmain(){
    void*initial_brk=sbrk(0); printf("Initialbrk:%p\n",initial_brk);

    void*new_brk=(char*)initial_brk+4096; if(brk(new_brk)==0){
        printf("brkmovedto:%p\n",new_brk);
    }else{
        perror("brkfailed");
    }

    return0;
}

```

how to simulate this in a minimal kernel, or implement a custom version in a toy OS

Goal: Simulate `brk()` in a Toy Kernel We'll assume:

A simple memory manager is already in place.

Your kernel has basic syscall handling.

The heap is a contiguous segment that can grow upward.

Step 1: Define Heap Region in Your Kernel

In your process struct (e.g., `task_struct`), define heap bounds:

c code

```
struct task_struct { ...  
    void* heap_start; void* heap_end; // current end of heap (brk)  
    void* heap_max; // maximum allowed  
};
```

During process creation:

c code

```
process->heap_start = alloc_some_memory(); process->heap_end = process->heap_start;  
process->heap_max = process->heap_start + HEAP_LIMIT; Step 2: Add the sys_brk System Call
```

c

Copy code

```
int sys_brk(void* new_end) {  
    struct task_struct* current = get_current_task();  
  
    if (new_end == 0) {  
        return (int) current->heap_end; // return current break  
    }  
}
```

```

if((char*)new_end<(char*)current->heap_start||
   (char*)new_end>(char*)current->heap_max){ return-1;//invalidrange
}

```

```

current->heap_end=new_end; return(int)new_end;
}

```

Step3:Hook sys_brk into Your Syscall Table Add a syscall number for brk, like:

c code

```
#define SYS_BRK 45
```

In your syscall dispatcher:

c code

```

switch(syscall_num){ ...
    case SYS_BRK:
        ret=sys_brk((void*)arg0); break;
}

```

Step4: Test brk() in Userland (if you simulate user programs) c code

```

int main(){
    void*current=syscall(SYS_BRK,0); printf("Heap start:%p\n",current);

    void*new_end=current+0x1000; syscall(SYS_BRK,new_end);

    void*check=syscall(SYS_BRK,0); printf("Heap now:%p\n",check);
}

```

This is a simplified view,

Ref.

Google

TutorialsPoint – `brk()` System Call in Unix

YouTube – `sbrk/brk` System Calls and Optimistic Allocation Explained